
Population Modeling Should Be First-Order in LLM-Based Survey Simulation

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 LLM-based survey simulation is often treated as a problem of generating more
2 realistic synthetic respondents. **We argue instead that population modeling**
3 **should be a first-order design problem in LLM-based survey simulation.** A
4 survey measures a target population, an LLM reflects an implicit model population,
5 and external domain evidence provides an observed but biased mediator population.
6 Using a real world large scale Chinese consumer-health survey benchmark and
7 a domain-matched corpus of nearly half a million online posts from Sina Weibo
8 and Xiaohongshu, we compare semantic mapping and LLM simulation under
9 post-level and population-structured aggregation. Across both estimator families,
10 recovering candidate mediator units before aggregation improves alignment with
11 real survey distributions. These results suggest that survey-simulation research
12 should report not only the model, prompt, or estimator, but also which population
13 is being represented, what mediator evidence is used, and what population unit
14 supports inference.

15 1 Introduction

16 LLM-based survey simulation has recently emerged as a promising tool for approximating public
17 opinion, consumer preferences, and population-level behavior. Researchers have prompted large
18 language models (LLMs) to answer surveys as synthetic respondents, personas, or demographic
19 groups, and have shown that these simulations can sometimes reproduce aggregate patterns in human
20 survey data [1, 2, 3, 4]. At the same time, domain-specific external evidence offers a way to ground
21 these simulations in observed patterns of language, behavior, experience, and preference.

22 However, current survey-simulation pipelines face a population-representation problem. When an
23 LLM is asked to answer a survey, its responses are shaped by the model’s training corpus, alignment
24 behavior, and implicit assumptions about the target population. When external evidence such as
25 social media, online trace, or other domain-specific observational data is used, it is often treated as an
26 unstructured pool of observations, even though such evidence is typically biased, heterogeneous, and
27 organized around distinct communities, contexts, and use cases.

28 **In this position paper, we argue that population modeling should be treated as a first-order**
29 **problem in LLM-based survey simulation: its reliability depends not only on how survey**
30 **answers are generated or estimated, but also on which population those answers are intended**
31 **to represent and what population unit the estimator operates on.**

32 We support this argument using a real-world Chinese consumer health survey and a domain-matched
33 corpus of Sina Weibo¹ and Xiaohongshu² posts. We organize the design space along two dimensions:

¹<https://en.wikipedia.org/wiki/Weibo>

²<https://en.wikipedia.org/wiki/Xiaohongshu>

the response estimator, comparing semantic mapping with direct LLM simulation; and the population unit, comparing post-level inference with population-structured inference over recovered social-media groups. Across both response estimators, we find that aggregating through latent social-media groups improves alignment with real survey distributions. Finally, we conclude with a call for survey-simulation research to treat population modeling as a first-order design problem, rather than relying only on more realistic prompts or stronger language models.




Figure 1: Population modeling with an observed mediator population. Survey simulation requires relating the LLM’s implicit population to the real survey target population through an observed but biased mediator population. In our setting, this mediator is a domain-specific social-media corpus.

2 Related Work

LLM-based survey simulation sits at the intersection of computational social science, population modeling, and language technology. We organize prior work around three design dimensions that are central to our argument: the *response estimator*, or how methods convert evidence into survey answers; the *mediator population*, or what observed evidence, if any, bridges the LLM and the target population; and the *population unit*, or what unit is simulated, inferred, and aggregated to form the final distribution estimate. Different choices along these dimensions encode different assumptions about which population is being represented. Making these choices explicit is therefore necessary for diagnosing when and why survey simulation succeeds or fails.

Response estimators. The dominant approach treats the LLM itself as the response estimator. “Silicon samples” and simulated-human-subject paradigms condition models on demographic profiles, personas, or respondent descriptions, elicit answers, and aggregate them into survey distributions [2, 1]. Subsequent work shows both the promise and limits of this strategy: LLMs can sometimes reproduce aggregate human patterns, but they also reflect implicit demographic and cultural priors, flatten individual heterogeneity, and fail under some inferential evaluations [3, 4, 5, 6]. Critical studies further show that plausible aggregate outputs can mask failures in regression coefficients, response biases, and downstream social-science inference [7, 8, 9]. A complementary line of work replaces direct generation with semantic representation. Semantic Similarity Rating (SSR) maps free-text responses to answer options through embedding similarity to reference anchors, producing response distributions rather than single generated answers [10]. Related critiques show that even non-generative mappings can carry systematic calibration biases [11]. Together, this literature suggests that the response estimator is not a neutral implementation detail: it determines which model assumptions enter the estimated survey distribution.

Mediator populations. Early LLM survey simulation often places no external mediator between the model and the target population; the model’s training data, alignment, and default priors directly shape the response distribution. Durmus et al. [12], for example, characterize the LLM’s implicit population by comparing model outputs with surveys from 65 countries and find that default responses are closest to U.S. and Western European opinion. This motivates methods that introduce external evidence to bridge the model and the target population. One established approach uses demographic

or census information as a structured mediator: multilevel regression with post-stratification corrects non-representative samples using population cell counts [13, 14], and recent LLM work extends this logic by constructing population-aligned personas or using rich individual interviews as grounding evidence [15, 6]. Other work uses domain-specific online evidence as an observed but biased mediator population. Such evidence can contain survey-relevant signals, but it is also shaped by selection effects, platform dynamics, and missing subgroups [16, 17]. Recent systems prompt LLMs to extract opinions from online text or use agentic pipelines to gather and aggregate such signals [18, 19]. However, these approaches often treat the mediator either as a flat pool of text or as a source requiring external demographic correction, leaving its internal structure under-modeled.

Population units. Every survey-simulation pipeline also makes an implicit choice about the unit of inference. Persona-based methods typically operate at the individual level: one profile, one simulated respondent, one response [2, 1, 5]. When richer individual evidence is available, such as interview transcripts, individual agents can become substantially more accurate [6]; when demographic structure is known, methods can instead aggregate through strata or persona mixtures weighted to match a target population [13, 14, 15, 20]. A shared assumption across much of this work is that the relevant population structure is known in advance or can be defined by demographic categories. In many domains, however, the available mediator evidence is organized around communities, topics, use contexts, or behavioral situations rather than clean demographic cells. Our work therefore treats the population unit as an empirical design choice: we recover candidate mediator units from observed domain discourse and test whether aggregating through these units changes alignment with real survey distributions. This supports a diagnostic question that prior work rarely isolates: holding the response estimator fixed, does the population unit change the estimated distribution?

3 Problem Formulation and Design Space

We study survey distribution estimation with three populations in mind: the *real survey population*, whose answer distribution is the target; the *LLM implicit population*, induced by the model’s training data and alignment; and an *mediator population*, which is biased but domain-relevant. Our goal is to use the observed mediator population as a mediator for estimating the real survey answer distribution.

For a survey question q with answer options $\mathcal{Y} = \{1, \dots, K\}$, let $p^*(y | q)$ denote the real survey distribution. Given a social-media corpus $\mathcal{S} = \{s_i\}_{i=1}^N$, we seek an estimator

$$\hat{p}(y | q, \mathcal{S}) \approx p^*(y | q).$$

Given a mediator population, we organize the method space along two dimensions, shown in Figure 2. The first dimension is the *response estimator*: how textual evidence is converted into a distribution over survey answers. The second dimension is the *population unit*: the unit over which response estimates are produced and aggregated.

3.1 Response estimator

A response estimator maps textual evidence x to an answer distribution $f(y | q, x)$. We consider two estimator families. Semantic mapping estimates f by comparing the semantic representation of x with answer-option anchors. LLM simulation estimates f by prompting a language model to directly return a probability distribution over the answer options. The detailed implementations are described in Sections 4.5.1 and 4.5.2.

3.2 Population unit

Under *post-level inference*, each social-media post is treated as an independent unit:

$$\hat{p}_{\text{post}}(y | q, \mathcal{S}) = \frac{1}{N} \sum_{i=1}^N f(y | q, s_i).$$

This treats the observed corpus as a flat collection of posts.

Under *population-structured inference*, the corpus is first organized into latent groups $\mathcal{G} = \{G_1, \dots, G_C\}$ with weights π_c such that $\sum_{c=1}^C \pi_c = 1$. The final estimate aggregates group-level




Figure 2: Two-by-two design space for LLM-based survey simulation.

113 response distributions:

$$\hat{p}_{\text{struct}}(y \mid q, \mathcal{S}) = \sum_{c=1}^C \pi_c \hat{p}_c(y \mid q, G_c).$$

114 This treats social media as a structured mediator population rather than a flat sample.

115 Together, these dimensions define four practices: post-level semantic mapping, population-structured
 116 semantic mapping, post-level LLM simulation, and population-structured LLM simulation. This
 117 design lets us separate errors due to response estimation from errors due to population representation.

118 4 Experiment

119 We use the experiment to support the central position of this paper: LLM-based survey simulation
 120 should be evaluated as a population-bridging problem, not merely as a response-generation problem.
 121 Accordingly, our empirical study is organized around three questions. First, what real survey
 122 population are we trying to approximate? Second, what observed social-media population can
 123 mediate between the LLM and that survey population? Third, does changing the population unit alter
 124 the estimated survey distribution?

125 4.1 The Target: A Real Survey Population

126 The target of estimation is a real-world Chinese consumer health survey with approximately 138,000
 127 responses. As summarized in Table 4 in Appendix A, the survey contains two types of questions:
 128 product-experience-oriented questions and demographic/background questions. Our main distribu-
 129 tional evaluation uses the six human-authored, single-choice product-experience questions in Panel A.
 130 The demographic and background questions in Panel B are used to characterize the survey population
 131 and support population-level analyses.

132 We clean the survey data by parsing each respondent’s questionnaire record from a JSON-encoded
 133 content field and repairing common formatting artifacts when necessary. We restrict the data to
 134 the target survey and aggregate valid responses into option-level answer counts for each question.
 135 Demographic and background questions are excluded from the main evaluation, and the final bench-
 136 mark contains six product-experience questions with empirical answer distributions from real survey
 137 responses.

4.2 The Mediator: An Observed Social-Media Population

The mediator population is a corpus of posts from Sina Weibo and Xiaohongshu in the Chinese consumer health / Traditional Chinese Medicine (TCM) domain. We collect approximately 448,832 posts focused on discussions of over-the-counter TCM products, including Huoxiang Zhengqi formulations, cough syrups, and cold medicines, primarily from a major pharmaceutical brand.

For the social-media corpus, we retain only Sina Weibo and Xiaohongshu posts that mention the target product family. We normalize text fields, remove hashtag wrappers, and filter out posts with confounding product meanings, recipes, parenting contexts, social mentions, bare reposts, or likely professional medical and health-care authors. We then remove exact duplicates and near-duplicates using a 10-character sliding-window deduplication rule to reduce template reposts and minor variants.

After filtering and deduplication, the corpus contains 27,838 posts. Following HDBSCAN clustering as detailed in Section 4.4, posts assigned to the noise label are removed, yielding the final 11,873-post corpus, organized into 66 recovered clusters. This approximately 12K-post corpus is used as the observed social-media population throughout the experiments.

4.3 The Test: Does the Population Unit Matter?

To test whether the population unit matters, we instantiate the four cells of the design space in Section 3. The comparison crosses two response estimators with two population units. The response estimators are semantic mapping and Qwen3-8B-based LLM simulation [21]. The population units are individual posts and HDBSCAN-recovered social-media groups: **(i) Post-level semantic mapping.** Apply semantic mapping to each post and average the resulting distributions; **(ii) Population-structured semantic mapping.** Apply semantic mapping within each recovered group and aggregate group-level distributions using cluster weights; **(iii) Post-level LLM simulation.** Prompt Qwen3-8B with a single post and aggregate the resulting post-level distributions; and **(iv) Population-structured LLM simulation.** Prompt Qwen3-8B with cluster-level context and aggregate the resulting group-level distributions using cluster weights.

The following subsections describe how we recover the candidate population units and how we instantiate the two response estimators.

4.4 Recovering Candidate Population Units

We recover latent social-media groups by clustering post embeddings. After product-name masking, text normalization, filtering, and deduplication, each post is encoded using bge-small-zh-v1.5 [22] for clustering. We then reduce the post embeddings to a 50-dimensional space using UMAP and apply HDBSCAN to identify dense semantic groups while treating low-density posts as noise. Posts assigned to HDBSCAN’s noise label are removed, yielding 11,873 posts organized into 66 clusters.

Representative recovered clusters are shown in Table 6 in Appendix B. The clusters capture distinct survey-relevant contexts, including medical use, heatstroke prevention, and product perception. This illustrates why we treat the social-media corpus as a structured mediator population rather than a flat collection of posts: different groups encode different use cases, attitudes, and consumption contexts that may map differently to survey answer distributions.

We do not claim that these clusters are demographic strata or that they recover the true survey population. Instead, they are candidate population units recovered from observed domain discourse. For each survey question, we select the Top- K clusters most relevant to the question before response estimation. Specifically, each cluster is represented by its topic description, and clusters are ranked by semantic similarity to the survey question. The Top- K clusters are then used as the mediator units for population-structured inference, with their within-cluster response distributions aggregated using the corresponding cluster weights.

Because the number of selected clusters can affect the trade-off between coverage and noise, we later conduct a Top- K sensitivity analysis (see Section 5.3), varying K and showing that performance is strongest in a moderate range, around $K = 10$ –15. Importantly, both clustering and Top- K cluster selection use only social-media text and survey-question text, and do not use survey responses. Thus, the population recovery and retrieval steps are unsupervised with respect to the evaluation target.

188 4.5 Response Estimators

189 To support our call that population modeling should be treated as a first-order design problem, we
190 instantiate both major estimator families in LLM-based survey research: non-generative semantic
191 mapping and generative LLM simulation.

192 4.5.1 Response Estimator I: Semantic Mapping

193 Semantic mapping provides a non-generative estimator of survey answer distributions. For each
194 survey question, we first convert every answer option into a natural-language anchor sentence. The
195 anchors are generated by Qwen3-8B using a fixed prompt that asks the model to rewrite each option
196 as a concrete sentence expressing the corresponding attitude or choice, while preserving the original
197 option order. This step produces one anchor sentence for each answer option.

198 Each anchor sentence and each social-media text unit are then encoded using bge-base-zh-v1.5 [23].
199 Given a text unit x and the anchor sentence a_k for option k , we compute their cosine similarity:

$$z_k(x) = \cos(e(x), e(a_k)).$$

200 The similarity scores are converted into a probability mass function over answer options:

$$f_{\text{sem}}(y = k \mid q, x) = \frac{\exp(z_k(x)/\tau)}{\sum_{j=1}^K \exp(z_j(x)/\tau)},$$

201 where τ is a temperature parameter controlling the sharpness of the similarity-based distribution.

202 For post-level semantic mapping, x is an individual social-media post. We compute one probability
203 vector for each post and average the resulting vectors across posts to obtain the final survey-distribution
204 estimate.

205 For population-structured semantic mapping, semantic mapping is first applied within recovered
206 social-media groups. For each cluster, we compute the semantic-mapping distributions of posts
207 associated with that cluster and average them to obtain a cluster-level distribution. The final estimate
208 is then obtained by aggregating cluster-level distributions using the corresponding cluster weights.
209 Thus, semantic mapping lets us test whether the same response estimator improves when it operates
210 over recovered population units rather than isolated posts.

211 4.5.2 Response Estimator II: LLM Simulation

212 LLM simulation provides a generative estimator of survey answer distributions. For LLM-based
213 response estimation, we use Qwen3-8B as a prompted distribution estimator. Rather than asking
214 the model to select a single option, we ask it to return a probability mass function over the answer
215 options in JSON format, enabling direct comparison with the empirical survey distribution.

216 For population-structured LLM simulation, the prompt conditions on a recovered cluster. It includes
217 the cluster topic, sampled posts from the cluster, the survey question, and the numbered answer-option
218 list. The translated prompt template is shown in Prompt 1 (Appendix C). The parsed output is a
219 vector $\{\text{"distribution"} : [p_1, p_2, \dots, p_K]\}$, where each p_k corresponds to one answer option.

220 For post-level LLM simulation, we use the same output format but condition the model on a single
221 social-media post. This prompt, shown in Prompt 2 (Appendix C), is used for the post-level LLM
222 corner of the 2×2 comparison.

223 The resulting post-level probability vectors are averaged according to the corresponding aggregation
224 rule. This provides a direct comparison between using the LLM at the individual-post level and using
225 it at the recovered-population-group level. In this way, the LLM is not only tested as a respondent
226 simulator, but also as an interpreter of candidate population units.

227 5 Evidence: Population Structure Changes the Estimate

228 We now use the 2×2 comparison to evaluate the central claim of this paper: LLM-based survey
229 simulation is not only a problem of response estimation, but also a problem of population representa-
230 tion. The results provide three pieces of evidence. First, structured social-media evidence improves

estimation across both response estimators. Second, this improvement is not explained by a single favorable cluster-weighting scheme. Third, the number of recovered groups matters in a way that is consistent with the population-bridging view: too few groups under-cover the mediator population, while too many groups introduce noisy tail contexts.

5.1 Evidence 1: Structured Social Media Outperforms Post-Level Evidence

Table 1 reports the aggregated comparison over the survey question set. The rows compare two response estimators: semantic mapping and LLM simulation. The first column reports a question-only baseline, where the estimator uses the survey question and answer options without conditioning on any external mediator population. The remaining two columns compare aggregation units when mediator evidence is used: post-level aggregation and population-structured aggregation.

For semantic mapping, the question-only baseline follows the standard SSR logic: each answer option is expressed as a reference anchor statement, and the survey question text is compared with these anchors to obtain an answer distribution. For LLM simulation, the model is prompted with only the survey question and answer options and is asked to output an estimated answer distribution directly. In both cases, the estimator operates without conditioning on any external mediator population.

The key pattern is that, once external mediator evidence is introduced, the population-structured aggregation cells perform best. For semantic mapping, moving from the question-only baseline to post-level mediator evidence reduces mean JS divergence from 0.0430 to 0.0295. Aggregating the same mediator evidence through recovered population units further reduces mean JS divergence to 0.0269. For LLM simulation, the question-only baseline performs poorly, with a mean JS divergence of 0.1581. Conditioning the model on individual posts reduces this value to 0.0328, while conditioning it on population-structured evidence further reduces it to 0.0275.

This supports our position in two ways. First, social media is useful as domain evidence, but it is more useful when treated as a structured mediator population rather than as a flat collection of posts. Second, the improvement appears across both semantic mapping and LLM simulation. Thus, the main lesson is not simply that one response estimator is better than another. Rather, the population unit changes the estimate.

Response estimator	Aggregation unit		
	Without social media	Post-level social media	Structured social media
Semantic mapping	0.0430	0.0295	0.0269
LLM simulation	0.1581	0.0328	0.0275

Table 1: Aggregated comparison of response estimators and population evidence. Values are mean Jensen-Shannon divergence over the question set; lower is better.

5.2 Evidence 2: The Gain Is Not Just a Weighting Artifact

A possible objection is that population-structured inference only works because the default cluster weights happen to be favorable. In our default aggregation, each recovered cluster is weighted by its empirical mass, i.e., the fraction of social-media posts assigned to that cluster. However, social-media mass is not necessarily survey-population mass: some contexts may be over-represented online, while others may be under-represented.

To test whether the result depends on this weighting choice, we keep the recovered clusters and the within-cluster predictions fixed, and vary only the cluster weights used in the final aggregation. Besides empirical cluster mass, we evaluate demographic reweighting based on province, age, and province-age combinations, as well as optimal-transport-based reweighting variants.

Table 2 shows that performance changes little across these choices. The default empirical weighting obtains a mean JS divergence of 0.0269, and all alternative weighting schemes remain within 0.0269–0.0272. This suggests that the benefit of population-structured inference is not driven by one particular weighting rule. Rather, the recovered groups themselves provide a stable intermediate representation for aggregating social-media evidence.

Cluster weighting scheme	Mean JS ↓
Empirical cluster mass	0.0269
Importance sampling by province	0.0271
Importance sampling by age	0.0269
Importance sampling by province + age	0.0271
Optimal transport by province	0.0272
Sinkhorn optimal transport by province	0.0269

Table 2: Robustness of population-structured semantic mapping across cluster weighting schemes.

5.3 Evidence 3: Population Structure Requires the Right Level of Granularity

Population-structured inference is not a claim that more groups are always better. It requires choosing a useful level of granularity. Using too few clusters may miss relevant subpopulations in the social-media corpus, while using too many clusters may introduce noisy or weakly related tail clusters. We therefore vary the number of retrieved clusters, K , and evaluate both population-structured semantic mapping and population-structured LLM simulation.

Table 3 reports the sensitivity results. For population-structured semantic mapping, performance is stable across K , with the best score at $K = 10$. For population-structured LLM simulation, performance is worse at $K = 5$, improves substantially at $K = 10$ and $K = 15$, and slightly degrades again at $K = 20$. Overall, the best region is around $K = 10$ – 15 .

This pattern is consistent with the population-bridging view. When too few clusters are used, the mediator population is under-covered: important contexts in the social-media corpus may be omitted. When too many clusters are used, weak or noisy tail groups can dilute the signal. The goal is therefore not to maximize the number of recovered groups, but to recover a level of population structure that is detailed enough to capture heterogeneity and stable enough to support aggregation.

Top- K	Population-structured semantic	Population-structured LLM
5	0.0270	0.0306
10	0.0269	0.0275
15	0.0275	0.0250
20	0.0277	0.0265

Table 3: Top- K sensitivity. Values are mean JS divergence over the question set.

6 Discussion: Implications for Future Survey Simulation

Our results support a broader position: future work on LLM-based survey simulation should treat population modeling as a first-order research problem. A real survey, an LLM, and a social-media corpus do not represent the same population. A survey measures a target population; an LLM reflects an implicit model population; and social media provides an observed but biased domain population. The central challenge is therefore not only to generate plausible survey answers, but to decide which population those answers are meant to represent.

This has implications for how the community should use social media. Social media should not be treated as a representative survey sample, nor should it be reduced to a flat pool of text. Its value lies in its observed, domain-specific structure: users discuss products, symptoms, preferences, routines, and use contexts in ways that may be relevant to survey answers. But this structure must be recovered and modeled explicitly. Otherwise, social-media-based survey simulation risks replacing one opaque population prior (the LLM’s implicit population) with another (the unexamined platform population).

This perspective also changes the role of LLMs. Much existing work asks whether LLMs can simulate individual respondents. Our results suggest that this may be the wrong default unit. When the available evidence is organized around communities, topics, or use contexts, the model may be more useful as an interpreter of recovered population units than as a generator of isolated synthetic individuals. A symptom-focused community, a seasonal consumption context, or a product-use scenario can provide a more coherent basis for distributional estimation than a single post or a lightly

specified persona. The broader design principle is that LLMs should be prompted at the level where evidence meaningfully represents population variation.

We therefore call for future work to make population assumptions explicit. Survey simulation papers should report not only what model they use and how they prompt it, but also what population mediates the estimate, what unit is aggregated, and why that unit is appropriate for the target survey distribution. Without this shift, LLM-based survey simulation may continue to produce plausible-looking answers whose population meaning remains unclear. With it, survey simulation can move toward more interpretable, auditable, and empirically grounded estimates of population-level opinion and behavior.

7 Counterarguments

Respondent simulation can be appropriate when respondent-level structure is informative. Our argument is not that personas or synthetic respondents are the wrong unit in general. When reliable individual-level covariates such as age, region, income, product usage, or prior attitudes are available and survey-relevant, respondent-level simulation can be a valid population-modeling strategy. In such cases, the respondent is a meaningful unit for representing population variation, especially when simulated responses are aggregated with appropriate population weights.

The key point is that respondent-level simulation should be justified rather than assumed. A persona is useful only insofar as it captures variation that matters for the target survey distribution. If the available covariates are incomplete, weakly related to the outcome, or mismatched to the target population, then simulating many respondents may create the appearance of coverage while still reproducing the LLM’s implicit population. Our position is therefore not anti-persona; personas are one possible population unit that should be evaluated against alternatives.

Recovered clusters are discourse units, not demographic populations. A possible concern is that clusters recovered from social-media text may represent topics, situations, or use contexts rather than demographic strata. This is exactly how we interpret them. We do not claim that unsupervised text clustering recovers demographic groups or the true survey population. Instead, the clusters organize heterogeneous mediator evidence before survey-distribution estimation. When individual-level covariates are unavailable, noisy, or incomplete, survey-relevant signals in social media may appear through product-use contexts, symptom narratives, seasonal routines, purchase situations, or shared attitudes. These clusters should therefore be understood as discourse-level mediator units: recurring contexts in the observed corpus, not demographic subpopulations.

This distinction is central to our position. Population-structured inference should not be read as claiming that topic clusters discover the true population structure. Rather, it tests whether the mediator corpus should be treated as a flat pool of posts or organized through recovered discourse units. If these units capture survey-relevant heterogeneity, aggregating through them should improve alignment with the target survey distribution; if they capture superficial or irrelevant variation, the comparison should show little benefit or even degradation. The value of the framework is not to privilege topic clusters over respondents or demographic cells, but to make the aggregation unit explicit evaluable.

8 Conclusion

In this position paper, we argued that the population unit should be a first-order design choice in survey simulation. Using Chinese consumer health surveys and Sina Weibo/Xiaohongshu posts, we showed that social media is more useful when its latent structure is recovered before aggregation. Across both semantic mapping and LLM simulation, population-structured inference outperforms post-level inference. This suggests that improving the response estimator alone is not enough; researchers must also ask what population unit the estimator operates on.

More broadly, we call for future survey-simulation work to make its population assumptions explicit. Researchers should report not only which LLM, prompt, or estimator they use, but also what mediator population is used, what unit is aggregated, and why that unit is appropriate for the target survey distribution. Progress in this area will depend on better methods for recovering, validating, and weighting mediator populations. Without such population modeling, LLM-based survey simulation risks producing realistic-looking answers with ambiguous population meaning. With it, survey simulation can become more interpretable, auditable, and empirically grounded.

References

- [1] Gati Aher, Rosa I. Arriaga, and Adam Tauman Kalai. Using large language models to simulate multiple humans and replicate human subject studies. In *Proceedings of the 40th International Conference on Machine Learning (ICML)*, pages 337–371. PMLR, 2023.
- [2] Lisa P. Argyle, Ethan C. Busby, Nancy Fulda, Joshua Gubler, Christopher Rytting, and David Wingate. Out of one, many: Using language models to simulate human samples. *Political Analysis*, 31(3):337–351, 2023.
- [3] Shibani Santurkar, Esin Durmus, Faisal Ladd, Cino Lee, Percy Liang, and Tatsunori Hashimoto. Whose opinions do language models reflect? *arXiv preprint arXiv:2303.17548*, 2023.
- [4] Jingwei Yang et al. Specializing large language models to simulate survey response distributions for global populations. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2024. University of Tübingen / University of Copenhagen.
- [5] James Brand, Ayelet Israeli, and Donald Ngwe. Using LLMs for market research. Technical Report 23-062, Harvard Business School, 2023. Available at SSRN: 4395751.
- [6] Joon Sung Park, Carolyn Popowski, Carrie Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. Generative agent simulations of 1,000 people. *arXiv preprint arXiv:2411.10109*, 2024.
- [7] James Bisbee, Joshua D. Clinton, Cassy Dorff, Brenton Kenkel, and Jennifer Larson. Synthetic replacements for human survey data? the perils of large language models. *Political Analysis*, pages 401–416, 2024.
- [8] Lindia Tjuatja, Valerie Chen, Tongshuang Wu, Ameet Talwalkar, and Graham Neubig. Do LLMs exhibit human-like response biases? a case study in survey design. *Transactions of the Association for Computational Linguistics*, 12:1011–1026, 2024.
- [9] Caleb Ziems, William Held, Omar Shaikh, Jiaao Chen, Zhehao Zhang, and Diyi Yang. Can large language models transform computational social science? *Computational Linguistics*, 50(1):237–291, 2024.
- [10] Benedikt Maier, Rishabh Jain, Simon Linke, et al. LLMs reproduce human purchase intent via semantic similarity elicitation of Likert ratings. *arXiv preprint arXiv:2502.XXXXX*, 2026. PyMC Labs / Colgate-Palmolive.
- [11] Eduardo Vera Pichardo. Measuring self-rating bias in LLM-generated survey data: A semantic similarity framework for independent scale mapping. *arXiv preprint arXiv:2602.13862*, 2026.
- [12] Esin Durmus, Karina Nguyen, Thomas I. Liao, Nicholas Schiefer, Amanda Askell, Anton Bakhtin, Carol Chen, Zac Hatfield-Dodds, Danny Hernandez, Nicholas Joseph, Liane Lovitt, Sam McCandlish, Orowa Sikder, Alex Tamkin, Janel Thamkul, Jared Kaplan, Jack Clark, and Deep Ganguli. Towards measuring the representation of subjective global opinions in language models. *arXiv preprint arXiv:2306.16388*, 2023.
- [13] Andrew Gelman and Thomas C. Little. Poststratification into many categories using hierarchical logistic regression. *Survey Methodology*, 23(2):127–135, 1997.
- [14] Wei Wang, David Rothschild, Sharad Goel, and Andrew Gelman. Forecasting elections with non-representative polls. *International Journal of Forecasting*, 31(3):980–991, 2015.
- [15] Eunbi Hwang et al. Alignedpersona: Population-aligned persona generation. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024. HKUST / Microsoft Research Asia.
- [16] Fernando Diaz, Michael Gamon, Jake M. Hofman, Emre Kiciman, and David Rothschild. Online and social media data as an imperfect continuous panel survey. *PLOS ONE*, 11(1):e0145406, 2016.

- 404 [17] Indira Sen, Fabian Flöck, Katrin Weller, Bernd Weiß, and Claudia Wagner. Correcting sociode-
 405 mographic selection biases for population prediction from social media. In *Proceedings of the*
 406 *11th International AAAI Conference on Web and Social Media (ICWSM)*, pages 628–631. AAAI
 407 Press, 2017.
- 408 [18] Luca Cerina and Raymond M. Duch. Artificially intelligent opinion polling: Using language
 409 models to analyse political opinions on social media. *arXiv preprint arXiv:2401.XXXXX*, 2024.
 410 University of Amsterdam / University of Oxford.
- 411 [19] Anonymous. SurveyPilot: Collecting human opinions from social media with agentic language
 412 models. In *Submitted to the Annual Meeting of the Association for Computational Linguistics*
 413 *(ACL)*, 2025.
- 414 [20] Bowen Wang et al. Mixture-of-personas language models for diverse opinion simulation. In
 415 *Findings of the Association for Computational Linguistics*, 2024. Yale University / CUHK.
- 416 [21] An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu,
 417 Chang Gao, Chengen Huang, Chenxu Lv, et al. Qwen3 technical report. *arXiv preprint*
 418 *arXiv:2505.09388*, 2025.
- 419 [22] Beijing Academy of Artificial Intelligence. BAAI/bge-small-zh-v1.5. <https://huggingface.co/BAAI/bge-small-zh-v1.5>, 2023. Hugging Face model card. Accessed: 2026-05-03.
- 420
 421 [23] Beijing Academy of Artificial Intelligence. BAAI/bge-base-zh-v1.5. <https://huggingface.co/BAAI/bge-base-zh-v1.5>, 2023. Hugging Face model card. Accessed: 2026-05-03.
 422

Question ID	Survey question	Option list
Panel A: Product-experience-oriented questions used for distributional evaluation		
Q1	How do you think the price of {Product Name}?	A: Very expensive; B: Somewhat expensive; C: Reasonable; D: Cheap
Q2	In what season would you be likely to buy {Product Name}?	A: Spring; B: Summer; C: Autumn; D: Winter; E: No seasonal preference
Q3	Which attribute matters when you buy {Product Name}?	A: Effectiveness; B: Safety; C: Brand reputation; D: Price; E: Others
Panel B: Demographic and background questions used for population characterization		
D1	Gender	A: Male; B: Female; C: Other / prefer not to say
D2	Age group	A: Under 18; B: 18–24; C: 25–34; D: 35–44; E: 45–54; F: 55 and above
D3	Province or city of residence	Free-text response

Table 4: Representative questions from the benchmark survey. Product names are masked as {Product Name}. Panel A shows product-experience-oriented questions used for distributional evaluation, while Panel B shows demographic and background questions used to characterize the survey population and support population-level analysis. Some questions and options are omitted for brevity.

User ID	Gender	Age	Location	Post time	Translated post
Panel A: Sina Weibo — short, conversational posts with incidental product mentions					
U-9100	Female	18–24	Guangdong	2023-07-21	“A summer without {Product Name} is not a proper summer.”
U-2775	Female	18–24	Zhejiang	2023-05-29	“Got serious heatstroke today—if it weren’t for {Product Name}, I think I would have passed out in the classroom.”
Panel B: Xiaohongshu — longer, structured health-diary posts with intentional product use					
U-4812	Female	34	Hubei	2023-10-08	“Used {Product Name} as part of a home-care routine ... described several ways of applying it, such as bath water or foot soaking ...”
U-4994	Female	25	Hunan	2023-09-09	“Shared a personal experience with {Product Name} ... emphasized moderate use, digestion-related feelings, and weight-management hashtags”

Table 5: Representative posts from the two social-media platforms. User IDs are anonymized, gender and age are obtained from publicly available social-media profile fields, locations are province-level, and post times are rounded to dates. Product names are masked as {Product Name}. Posts are translated and lightly paraphrased to preserve platform-level discourse patterns while minimizing privacy risk and avoiding disclosure of identifiable user-level information.

424 B Cluster Illustration

ID	Size	Theme	Cluster topic
Medical use			
56	1,518	GI / digestive	Users describe diarrhea, vomiting, abdominal pain, and other digestive symptoms, often presenting {Product Name} as relief.
53	886	Cold / flu / fever	Users discuss cold, fever, and gastrointestinal discomfort, linking the product to everyday symptom management.
Heatstroke prevention			
64	514	Heatstroke experience	Users share heatstroke experiences and describe using {Product Name} for immediate relief in outdoor settings.
62	320	Summer preparedness	Users discuss summer heat-prevention routines, including carrying {Product Name}, sunscreen, and water.
Taste and product perception			
55	525	Taste, negative	Users strongly dislike the taste while still acknowledging the product's therapeutic value.
35	181	Positive perception	Users express highly positive evaluations and describe the product as a "miracle drug" (<i>shényào</i>).

Table 6: Representative clusters from the recovered social-media population. Cluster sizes indicate the number of posts assigned to each recovered group after filtering, deduplication, clustering, and HDBSCAN noise removal.

425 C Prompt Templates

Prompt 1: Population-structured LLM simulation prompt

You are a consumer research analyst.
The following social-media posts are samples from a consumer group (topic: {topic}):
Posts: {posts}
Based on the consumer characteristics, use contexts, and attitudes reflected in these posts, estimate this group's choice proportions for each option in the following survey question.
Question: {question}
Options: {options_numbered}
Please strictly output the option proportions in JSON format. The values must be decimals between 0 and 1 and must sum to 1:
{"distribution": [proportion for option 1, proportion for option 2, ...]}
JSON:

Prompt 2: Post-level LLM simulation prompt

You are a consumer research analyst.
The following is a social-media post from a consumer:
Post: {post}
Based on the consumer characteristics, use context, and attitudes reflected in this post, estimate this consumer's choice proportions for each option in the following survey question.
Question: {question}
Options: {options_numbered}
Please strictly output the option proportions in JSON format. The values must be decimals between 0 and 1 and must sum to 1:
{"distribution": [proportion for option 1, proportion for option 2, ...]}
JSON:

NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes], [No], or [N/A].
- [N/A] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a proper justification is given (e.g., error bars are not reported because it would be too computationally expensive” or “we were unable to find the license for the dataset we used”). In general, answering [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- Delete this instruction block, but keep the section heading “NeurIPS Paper Checklist”.
- Keep the checklist subsection headings, questions/answers and guidelines below.
- Do not modify the questions and only use the provided macros for your answers.

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not include experiments.

- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).

- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The authors should answer **[Yes]** if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.

- 632 • The paper should disclose whether the full research project required more compute
633 than the experiments reported in the paper (e.g., preliminary or failed experiments that
634 didn't make it into the paper).

635 **9. Code of ethics**

636 Question: Does the research conducted in the paper conform, in every respect, with the
637 NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

638 Answer: **[TODO]**

639 Justification: **[TODO]**

640 Guidelines:

- 641 • The answer [N/A] means that the authors have not reviewed the NeurIPS Code of
642 Ethics.
643 • If the authors answer [No], they should explain the special circumstances that require a
644 deviation from the Code of Ethics.
645 • The authors should make sure to preserve anonymity (e.g., if there is a special consid-
646 eration due to laws or regulations in their jurisdiction).

647 **10. Broader impacts**

648 Question: Does the paper discuss both potential positive societal impacts and negative
649 societal impacts of the work performed?

650 Answer: **[TODO]**

651 Justification: **[TODO]**

652 Guidelines:

- 653 • The answer [N/A] means that there is no societal impact of the work performed.
654 • If the authors answer [N/A] or [No], they should explain why their work has no societal
655 impact or why the paper does not address societal impact.
656 • Examples of negative societal impacts include potential malicious or unintended uses
657 (e.g., disinformation, generating fake profiles, surveillance), fairness considerations
658 (e.g., deployment of technologies that could make decisions that unfairly impact specific
659 groups), privacy considerations, and security considerations.
660 • The conference expects that many papers will be foundational research and not tied
661 to particular applications, let alone deployments. However, if there is a direct path to
662 any negative applications, the authors should point it out. For example, it is legitimate
663 to point out that an improvement in the quality of generative models could be used to
664 generate Deepfakes for disinformation. On the other hand, it is not needed to point out
665 that a generic algorithm for optimizing neural networks could enable people to train
666 models that generate Deepfakes faster.
667 • The authors should consider possible harms that could arise when the technology is
668 being used as intended and functioning correctly, harms that could arise when the
669 technology is being used as intended but gives incorrect results, and harms following
670 from (intentional or unintentional) misuse of the technology.
671 • If there are negative societal impacts, the authors could also discuss possible mitigation
672 strategies (e.g., gated release of models, providing defenses in addition to attacks,
673 mechanisms for monitoring misuse, mechanisms to monitor how a system learns from
674 feedback over time, improving the efficiency and accessibility of ML).

675 **11. Safeguards**

676 Question: Does the paper describe safeguards that have been put in place for responsible
677 release of data or models that have a high risk for misuse (e.g., pre-trained language models,
678 image generators, or scraped datasets)?

679 Answer: **[TODO]**

680 Justification: **[TODO]**

681 Guidelines:

- 682 • The answer [N/A] means that the paper poses no such risks.

- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.